



SHIKHAR GHOSH
ZOE CULLEN
SHWETA BAGAI

OneDigital (B): Building a Digital Workforce

If we're capable of replicating our employees and their cognition into an artificial system, how do we then balance that with valuing human capital and human output?

— Vinay Gidwaney, Chief Product Officer, OneDigital¹

Introduction

In late 2025, Vinay Gidwaney (Chief Product Officer, OneDigital) reviewed the results of the AI strategy he had presented to the CEO earlier in the year. Gidwaney had framed OneDigital not as an insurance or technology company, but as a human capital-driven advice firm delivering consulting services across employee benefits, retirement and wealth management. Their product was advice, historically generated by experienced producers and consultants whose judgment was shaped by years of domain expertise and client interaction. AI challenged this premise directly. The technology could incorporate the expertise of OneDigital's best producers and make that available to every client at a fraction of the cost. AI had also severed the linear relationship between producers and revenue that had defined OneDigital's growth model for decades.

Gidwaney had spent months working through these tensions. He built a team, hired AI specialists, ran experiments and developed what he believed was a differentiated approach. The company would pursue Cognition as a Service (CaaS)—equipping employees with digital coworkers that provided access to unlimited brainpower. OneDigital would focus on augmenting rather than replacing human capability. They would treat AI systems as employees rather than software tools; AI deployment would be an HR transformation project instead of a technology initiative. The team recognized that the hardest challenges were about people and culture and not algorithms.

By mid-2025, the company had deployed a dozen AI coworkers handling over 200,000 employee conversations. Each AI coworker was developed using the same methodology: identification of a business need, creation of a job description, selection of a supervisor, progression through intern and apprentice stages, and finally a promotion to full availability. AI coworkers were trained by humans

¹ All quotes from Vinay Gidwaney are sourced from casewriters' interviews with him in January 2026.

with tightly controlled source documents. They were conversational and could generate content for any audience based on the specific need of the OneDigital employee. Immediate results were encouraging with an adoption rate of 85.5% in deployed units. Producers using AI coworkers were servicing their clients better and had measurable impacts on revenue and retention. In 2025 alone, AI tools returned 59,165 hours to the workforce—capacity redirected toward relationship-building and revenue growth.

If AI could replicate expertise, how should the company value human talent? OneDigital had built a 25-year acquisition strategy around a powerful story: “join our platform, serve your clients, and benefit from our scale and expertise.” What would happen to that story as the company extracted cognitive capabilities from acquired firms and embedded them in AI systems that made human expertise less scarce? Moreover, was moving intelligence out of their employees and into systems the path to creating the company’s most valuable asset?

OneDigital’s AI Approach

In mid-2024, OneDigital's leadership identified three distinct ways that companies were incorporating AI into their operations. The first was transactional usage: employees asking the AI systems to do simple and independent tasks like answering factual questions, summarizing/editing documents or taking meeting notes. Most organizations that used AI remained at this level. The second approach was collaborative: employees engaging in extended conversations with AI systems, using them as thinking partners for complex problems. This usage pattern was common in personal contexts, where individuals might maintain ongoing dialogues with AI about parenting, travel planning or health decisions. This was rare in corporate settings due to concerns about confidentiality of corporate information. The third approach (which companies like Salesforce had begun discussing publicly) was to have AI manage corporate processes from end to end. The AI system could execute multi-step tasks such as screening job applicants, writing code for new software applications, generating marketing materials or managing customer interactions. Companies were still experimenting with the level of autonomy they were willing to give AI systems.

Collaborative AI

OneDigital made a strategic choice: the company would focus on the second approach of collaborative AI instead of pursuing full automation. This strategy represented a bet about the future of professional services. (See **Exhibit 1: OneDigital’s AI Journey**). Gidwaney explained the reasoning, “If we put AI between us and our clients, what’s our purpose? Small business owners don’t want to talk to a chatbot about their biggest challenges.” In an advice-giving business such as OneDigital, the human relationship remained the core asset. Clients might accept AI-generated recommendations, but OneDigital believed they wanted those recommendations delivered through trusted human advisors. The company’s competitive advantage was built on trusted relationships and judgment that emerged from understanding a client’s complete context. While AI could handle certain tasks, OneDigital concluded that client-facing relationships needed to remain human.

This view distinguished OneDigital from companies pursuing full automation. In some industries such as airlines, customers might not want to interact with humans and were keen to only get their problems solved efficiently. In such contexts, replacing humans with AI made economic sense. On the other hand, in complex B2B relationships involving significant financial decisions and ongoing advisory work, OneDigital believed human trust remained essential. The company’s AI strategy was therefore designed to increase its employees’ capabilities rather than replace them. The decision reflected a strategic bet that better-equipped consultants would win additional business and serve clients more effectively than pure AI-driven competitors.

The Engineering Transformation

In 2024, OneDigital implemented a hiring freeze for engineering positions. Engineers who left the company were not replaced. The remaining engineers were reassigned: they would no longer build features but instead maintain platforms and infrastructure. Product managers, embedded directly in business units, assumed responsibility for building features using AI coding assistants such as Replit and Claude Code. These tools allowed non-engineers to create software through natural language descriptions. Product managers were thus able to build custom applications in days rather than months. Gidwaney explained, “If we can deploy AI capabilities to allow product managers to build things quickly and on the fly—as connected to the use cases and the users as possible—then we can change the ROI calculation on how to build things.” With a view to maintain quality and safety standards, a few dozen approved individuals were given access to AI coding tools.

The engineers’ new role was to maintain the “vibe coding infrastructure”—the underlying platforms that handled production support, DevOps, database management, scalability, security, privacy and compliance. This infrastructure ensured that applications built by non-engineers met the company’s standards for reliability and data protection. A full-stack engineering team cost \$250-\$350,000 per year. Under the traditional model, only projects with million-dollar revenue potential could justify that investment. Other ideas applications that might save 100 employees a few hours per week or improve quality in ways that didn’t directly generate revenue remained unbuilt because they could not meet the ROI threshold. With AI coding platforms, these “long tail” ideas became feasible. Business units gained the ability to solve their own problems without competing for scarce engineering resources.

The HR-Centric Approach

Early in its AI planning, OneDigital determined that the most important department in a company when it came to AI was human resources. Gidwaney commented, “We are in the people business. We run the HR department for thousands of employers around the country. We knew from the get-go that we needed to take an HR-centric approach to how we think about AI.” The company recognized that AI would reshape fundamental aspects of work: performance evaluation, merit determination, skill valuation and team structure. These were HR and not IT challenges. Rather than treating AI as a back-office efficiency tool, Gidwaney viewed it as a form of talent that competed directly with human cognition. This framing became the guide for nearly every subsequent decision on implementation. (See **Exhibit 2: Five Dimensions of Workplace Transformation**).

Framing the AI Story

Recognizing that an AI strategy had to explicitly address employees’ concern that AI would take their jobs, the company invested heavily in communication. At OneDigital’s 2024 national conference, which brought together 1,500 employees, Gidwaney delivered a keynote address focused entirely on principles and outcomes (See **Exhibit 3: AI Principles**). Instead of dwelling on technology or implementation details, the presentation explained how AI would give employees time back and help them generate better insights for clients. The framing was explicit: AI tools were designed to make employees more effective, not to replace them. By addressing the job security concern directly and centering the narrative on employee empowerment, OneDigital reduced resistance to adoption considerably.

The company focused initial deployments on employees who were already high performers. These individuals were typically productivity-focused (always seeking better tools and more efficient workflows), revenue-oriented (understanding that their income depended on results) and comfortable with experimentation. For this group, AI represented an obvious enhancement to their

capabilities, and they adopted it with minimal resistance. The company gave these early adopters access to multiple AI tools (ChatGPT, Claude, Perplexity, and custom platforms). Gidwaney reflected, “The best way to gain value in your company is to find the A players in the business and give them these tools to become A+ players.” The only requirement was participation in weekly calls where users shared their learnings. This approach allowed OneDigital to discover actual use cases based on how people worked, rather than trying to mandate usage patterns from the top down. Gidwaney elaborated, “When my team introduces coworkers, the first thing that employees think is that this is designed to make my life easier. This is designed to help me serve my clients. Those are things that I really care about. They’re not thinking this is designed to take my job away. We have to set that premise really early on through that storytelling.”

Not everyone embraced the transition smoothly. The most consistent pushback came from senior producers who viewed their expertise as their “secret sauce” — accumulated over decades and deeply tied to client relationships. Gidwaney’s response was pragmatic rather than reassuring. “I’m not trying to say that what you have to offer is invaluable anymore. I’m just saying that it’s not a protectable thing that you think you have anymore.” He argued that the real threat was not OneDigital’s AI systems but competitors who would be “more data-driven, more proactive, more up-to-date” and could replace what they had to offer. “A good handful of folks just don’t believe in that,” Gidwaney admitted. “They will retire out of the business. And we’re okay with that.”

Strengthening the Organizational Muscle

The AI transformation was not going to be a one-time change but building what Gidwaney called “organizational muscle,” or the ability to blend technology thinking and business thinking continuously at every level. When an embedded product manager sat inside a practice, understood the business deeply enough to identify a need, built a working solution in days, deployed it on existing platforms and measured its impact—that was a business leader who thought in technology terms. The first time a practice would work this way, there was a learning curve. By the fifth iteration, practices could generate their own ideas and execute within sprints. He explained, “Our competitors will buy AI tools. They will hire consultants. They will run pilot programs. But they will not have spent years building the organizational muscle of integrating technology into the daily rhythm of how every practice operates. That instinct cannot be purchased or shortcut. It has to be built, repetition by repetition, win by win.”

Developing Digital Coworkers

The AI Coworker Team

In September 2023, OneDigital made its first hire specifically for AI implementation: Sean Daly, who had previously worked in data science and held degrees from Harvard and MIT. A year later, the company hired Robin Singhvi as Director of AI Programs. Singhvi had previously worked at Welltok and had experience as a startup entrepreneur. The team that Daly and Singhvi built was unusual for an AI initiative. Rather than focusing primarily on engineering and data science, the AI coworker team emphasized skills more commonly associated with recruiting and talent management. Their core responsibility was to staff the company with AI employees. Gidwaney explained, “They have more HR skills than technical skills and their job is to work with P&L leaders across the business.” The process worked as follows: The AI Programs team would meet with leaders from OneDigital’s various business units and discuss those leaders’ talent needs. What capabilities did they lack? What bottlenecks slowed down their teams? If they could hire anyone, what would that person do? Based on these conversations, the AI coworker team would create a job description. This was not a technical specification or a product requirements document. It was a standard job description that might be

posted on LinkedIn or Indeed, specifying responsibilities, required knowledge, communication style and expected outcomes.

The team then created a “resume” for a worker who met these requirements. The resume included experience (such as “worked in employee benefits consulting for eight years, specializing in high-cost claimant analysis”), personality traits (“thoughtful communicator who asks clarifying questions before making recommendations”) and even biographical details. Research had indicated that people related more naturally to AI systems that were individualized rather than generic. From a user's perspective, AI should feel like a colleague with a distinct personality and specific expertise. (See **Exhibit 4: Map AI Features to Human Qualities**).

From a technical standpoint, AI coworkers were not particularly sophisticated. Gidwaney explained, “Our goal is to build a methodology that encompasses what it takes for us to ‘hire,’ ‘onboard,’ and ‘manage’ these AI coworkers as they become effective members of our team.” The company developed a chatbot framework hosted on Microsoft Azure, using standard LLM APIs from OpenAI and Anthropic. The elements were all implemented as system prompts—instructions to the underlying large language model about how to respond. “Coworkers are there to help you think. Chatbot based environments are the best way for that to happen. We can have a conversation. We can go back and forth. You're challenging me. I'm challenging you. We're thinking better collectively.” The coworkers had access to curated knowledge bases (typically documents about OneDigital's services, industry regulations and best practices) and conversation history for context. Gidwaney described the technology stack as “pretty vanilla”—any company with similar resources could license comparable technology. He elaborated: “Building the coworker is not the hard part. It's really getting to the point of engaging all the stakeholders, writing all the prompts, making sure you've got all the knowledge base elements correct. There's no coding, there's no engineering, it's really about that product manager translating the idea into the prompting.”

AI Coworkers: Career Progression

OneDigital implemented a three-stage career progression for AI coworkers: intern, apprentice, and full-time employee. (See **Exhibit 5: The Coworker Hiring Process**). During the intern stage, the AI coworker team collaborated with business unit leaders to design and test the initial concept and build out their resume. Could an AI actually perform this function? Did it have access to the necessary knowledge? Could it communicate effectively with the intended users? What additional information, tools or capabilities did the system need to perform its function? Several AI interns failed this stage. Gidwaney explained, “We have dozens of AI coworkers that were interning that have been left on the cutting room floor because they just didn't hit the performance that we wanted to. We may go back and revisit them because LLMs continue to get better.”

Successful interns advanced to apprenticeships, a stage which typically lasted two to three weeks. At this time, the AI was made available to a carefully selected group of employees who would use it for its intended purpose. These pilot users provided detailed feedback: What worked well and where did the AI fall short? What improvements would make it more useful? The AI coworker team incorporated this feedback to refine the prompting. They adjusted its knowledge base, communication style and its ability to handle edge cases and unexpected questions, as needed. When an AI apprentice demonstrated consistent value and sufficient use cases were performing well, it was promoted to full-time availability. At this third stage, any employee in the company could access the AI through OneDigital's internal systems. Conversations were retained so users could maintain context across multiple interactions. The ability to share conversations with colleagues was built in, allowing teams to collaborate around the AI's recommendations (See **Exhibit 6: AI Coworker Evaluation and Progression**).

Promoting AI coworkers was different from promoting humans. With human employees, direct supervisors made promotion decisions, though they might harbor biases—wanting to retain top talent, avoid resource competition or allowing personal attachments to cloud judgment. For AI coworkers, supervisor bias could be worse—a high-performing AI might substitute for the supervisor's own work. Unlike humans, AI coworkers could undergo structured examinations and head-to-head model comparisons. Evaluations and testing allowed assessing whether an AI coworker was ready to move from an intern to apprentice to full-time availability.

The Supervisor Model

OneDigital's most significant organizational innovation was the requirement that every AI coworker have a human supervisor, and that this supervisor be in the business unit, not in the technology organization. Supervisors were given clear responsibilities: training the AI (writing and refining all system prompts), defining what knowledge the AI could access, reviewing the AI's outputs, and continuously improving performance based on user feedback and changing business needs. By mid-2025, OneDigital's AI coworkers had participated in more than 200,000 conversations with employees. Since it was not feasible for supervisors to review each interaction individually, other coworkers sampled/analyzed conversations and flagged potential issues for human review. Business unit leaders were accountable for performance of their teams, both human and AI. If an AI coworker provided poor advice or failed to meet expectations, the business leader had to address the issue.

The product team recommended a simple heuristic for selecting AI supervisors: identify the person who performed the relevant job best and make them the supervisor. This meant pulling top performers out of direct revenue-generating roles—a decision that required careful justification as these were the employees generating the most immediate value. The emphasis on having qualified supervisors raised several organizational issues. How much time and training would this require from some of the most valuable employees? Should the AI coworker's supervisor be a content expert or the manager of a unit? How should the company account for lost productivity in setting up AI coworkers?

The Case of Ben and Shelley

OneDigital created an AI coworker “Ben the Benefits Expert” to help consultants advise clients on managing health insurance costs. (See **Exhibit 7: Ben the Benefits Expert**). This was a pressing issue for nearly every employer given the industry's double-digit premium increases. Ben's job description stated that he could analyze an employer's health plan data, identify cost drivers and recommend strategies ranging from plan design changes, self-insurance options to pharmacy benefit optimization. His expertise emerged from extensive training in health plan structures, claims analysis and cost management techniques.

The product team recommended that Shelley McLean, a OneDigital top benefits consultant be Ben's supervisor. McLean had built a \$10 million book of business and consistently won new clients while maintaining high retention rates among existing ones. This created immediate tension: pulling a key revenue generator off active client service to train an AI system. Making McLean Ben's supervisor implied that she would be taking on a new role. She would also need to learn the craft of prompt and context engineering; conduct training sessions with Ben; hold office hours where consultants could work through complex cases; and continuously update Ben's knowledge base if regulations changed or the team learned new strategies.

The decision reflected OneDigital's broader philosophy about AI supervision. In McLean's case, she was encoding her expertise into a system that could then be deployed to hundreds of consultants.

Interestingly, McLean's role evolved in an unexpected direction. Rather than simply training Ben, she began using the AI coworker as her “megaphone”—a medium for transmitting best practices across OneDigital's entire employee benefits practice. When McLean conducted webinars or training sessions, she showed consultants how to access that knowledge through Ben. Gidwaney observed, “Ben became the medium by which training and knowledge transmission happened. It was a much better medium than videos or documents.”

Ben's impact extended beyond McLean's direct supervision. An HR Specialist used Ben to generate job descriptions, clarify confusing state regulations, and compare versions of employee handbooks. A Managing Consultant used Ben to consolidate and compare Third-Party Administrator options for a client. By uploading demonstration transcripts and asking for side-by-side comparisons, she offloaded 75% of her administrative work, allowing her to “focus on refining the deliverable and adding her own insights.” These employees weren't just using Ben; they were effectively accessing McLean's expertise, at scale. This also raised questions the company had not fully resolved: Should the business unit be given budget relief if McLean's book of business declined? Should her compensation reflect the impact of her work through Ben?

Use Cases and Impact

By mid-2025, OneDigital had deployed 11 coworkers in full-time service across every major practice—Employee Benefits, Property & Casualty, Financial Services, M&A and cross-practice function. (See **Exhibit 8: Examples of AI Coworkers**). Another seven were in the intern stage and six were progressing through apprenticeships. The company had also begun developing capabilities to integrate these AI systems into Microsoft Teams group chats, where they could be mentioned with the @ symbol and would participate in conversations like human team members. Each coworker addressed a specific business need identified by practice leaders, progressed through intern and apprentice stages, and was supervised by domain experts within business units. Human team members could “pick an AI coworkers off the shelf” based on their immediate need and have a collaborative engagement with it.

Usage data suggested that engagement was both frequent and sustained. An estimated 85.5% of users interacted with AI coworkers weekly or daily—more than four times the industry average of 15-20%. User satisfaction stood at 73%, with employees averaging 107 conversations per user per month. Majority of users reported increased productivity, with the most common applications involved client-facing support (e.g., answering client questions and providing service responses), summarizing documents and responding to client queries. AI coworkers were primarily augmenting human judgment and responsiveness, especially in advisory and service roles, rather than automating end-to-end processes. The productivity gains were substantial. A Senior Client Executive reported that his 10-person team was saving between 10 and 15 hours per week per employee—the equivalent of adding a full-time staff member. A Senior Account Manager described 30x time savings on a routine task. For example, searching for a missing answer on a policy form—a process that previously required calling multiple underwriters and could take over an hour—was completed in under two minutes by her AI coworker, which returned “not only the correct answer, but with the exact page number in the policy document.”

Beyond time savings, groups experienced improved morale. “People are less overwhelmed, more engaged, and have time to focus on strategic projects and client relationships. The quality of our work and our customer service has gone up.” Moreover, employees found diverse applications that the product team had not anticipated. A Senior Client Service Specialist used AI to explain commercial insurance concepts to Spanish-speaking prospects, work that “would have normally taken me an entire day to research, type up, then translate but was completed in a matter of five minutes.” Survey data also revealed that employees that had experienced what AI could do, wanted to go further. An

estimated 54% of employees wanted hands-on workshops, 50% wanted documentation and guides and 49% wanted video tutorials.

Financial Mandate and Future Architecture

The coworker portfolio demonstrated OneDigital's approach: rather than building general-purpose AI, the company created specialized coworkers mapped to specific workflows, supervised by practitioners who understood both the work and the clients. While OneDigital emphasized voluntary adoption and bottom-up experimentation, the company also implemented a top-down requirement that accelerated AI integration across the business. For the 2026 budget planning process, the CFO added a new requirement: every business unit leader had to address how they would use AI as part of their budget submission. Leaders who could not articulate an AI strategy faced additional scrutiny during budget reviews. For the 2027 planning cycle, this requirement would become more stringent: budget approval would be contingent on demonstrating an AI implementation plan. By making AI strategy a prerequisite for budget approval, the CFO and CEO signaled that AI integration was not optional. The mandate reflected broader pressure. During OneDigital's recent recapitalization discussions, potential PE investors had brought in their AI experts to evaluate opportunities for margin improvement and EBITDA growth, signaling that AI strategy would be central to any future partnership.

Overall, the company was playing a significant role in helping clients transition to a hybrid workforce (i.e., humans and AI working together). (See **Exhibit 9: A Hybrid Workforce**) The company was developing an AI product portfolio with distinct layers. The coworker layer was focused on conversational AI for thinking and collaboration (as described). The agent layer emphasized autonomous task completion and focused on AI systems that could execute multi-step processes without human intervention once given an instruction. The builder layer comprised AI coding platforms that allowed business users to create custom applications. An emerging fourth layer called Orchestration would coordinate across these capabilities—enabling coworkers to invoke agents, agents to request that builders create new tools and builders to generate new coworkers or agents (as needed).

The company was also developing more proactive and deeply integrated AI systems. The current generation of AI coworkers were reactive; they responded when prompted. The next generation would monitor workflows and suggest actions before being asked, join relevant conversations automatically, and anticipate needs based on context. Gidwaney projected, “Five years from now, how much of what we do with our clients is AI-originated? It will be 100%. Because every single recommendation, every piece of data, every service that we provide to our clients will originate from some AI brain saying that this is what their client should do.” This created the question: what would humans do?

Anticipating that the majority of work producers did daily was changing, the company started a program ODPP (OneDigital Producer Program), an effort to imagine the producer of the future. Gidwaney explained: “The producer of the future is really an AI orchestrator and relationship builder more than they are a benefits retirement expert,” drawing parallels to software engineering where engineers had shifted from hands-on coding to understanding business requirements and prompting and managing their agents rather than writing code. A coworker DEX, designed specifically for the ODPP was on track to reshape sales roles. All producers had to do was input a prospect's name; DEX would conduct deep research, perform ideal customer profile mapping, score the opportunity, identify key decision-makers, generate tailored outreach language and surface similar customers OneDigital already served. Gidwaney explained: “It will do all the work for you, and your job is to

read that email, copy/paste and send, and then pick up the phone and call them.” The skill map of what the company was looking for from their talent was going to change.

AI had also fundamentally altered OneDigital's acquisition strategy, their driver of inorganic growth. The company continued to pursue deals at a steady pace, but leadership increasingly evaluated targets through a new lens: cognitive assets that could be extracted and scaled. “We're buying a lot of specialty businesses,” Gidwaney explained. He described the recent acquisition of a firm specializing in student health plans. OneDigital would extract the expertise from the three-person team, encode it into an AI system and make it available to every office nationwide. This created tension as acquired firms expected their people and knowledge to remain central to client relationships. Instead, their expertise was being democratized—made accessible to thousands of OneDigital employees through AI, potentially eroding what had made the acquired firm valuable in the first place.

Broader Implications

By late 2025, OneDigital's AI coworker initiative had delivered measurable results and employees were embracing AI tools designed as collaborative partners. The company had achieved 85.5% weekly or daily AI usage—more than four times the industry average of 15-20%—with 73% user satisfaction. Despite the progress, Gidwaney found himself grappling with fundamental contradictions about value creation in an AI-augmented world.

Valuing Producers

OneDigital's financial model had always been straightforward: more producers meant more revenue, with predictable ratios. But with AI breaking the linear relationship between headcount and output, how should the company measure value? Gidwaney reflected:

If we hire 150 producers, it's going to result in \$300 million in revenue. With those 150 people, it's not just an issue of saying we only need 100 because we can make them more effective in their selling with AI. When I look at the value of 100 people and what they're capable of and I'm replacing that with AI, how do I quantify that value? Do I quantify it in revenue or some other measure of intelligence, cognition, and capability?

The answer would determine compensation structures, hiring strategies, and ultimately the kind of talent the firm valued. “It's not apples-to-apples math,” Gidwaney observed, comparing incentive-heavy producer compensation to the cost of deploying AI systems.

Equity and Access

Producers who mastered AI tools generated more revenue, creating a new dimension of inequality within the workforce. Simply providing access wasn't sufficient; some people could leverage AI in certain contexts far more effectively than others whether due to aptitude, training, or simply comfort with technology. Gidwaney described, “There's a jagged edge of AI within institutions themselves, within their own workforces that could lead to inequities, fairness and compensation issues that are not well understood.” If AI could compress the learning curve and amplify junior employees' capabilities, what would happen to the traditional value of experience? How do they provide equitable access to AI to allow all employees to be successful?

Defensibility

OneDigital had built its strategy around the belief that in complex advisory relationships, clients would continue to value human intermediaries even as AI capabilities improved. As AI systems became more sophisticated and capable of longer conversations and better judgments, would clients continue seeing value in the human layer? Or would the cost differential eventually overwhelm the trust premium? Gidwaney commented, “If AI is capable of generating that advice, then how do we shape our talent in the future so that there are things that humans can do that are unique to them. There are things that the AI can do that it’s well-positioned and equipped for. How do we quantify the value of that asset?”

The company's AI approach had implications for both cost structure and competitive positioning. OneDigital might not achieve the same cost reductions as competitors who replaced consultants with AI. Although OneDigital focused on augmentation, small upstarts in more transactional segments were experimenting with AI-first models that minimized human roles. If their bet was correct, OneDigital would differentiate on quality and trust. Private equity investors explicitly evaluated AI through the lens of margin improvement and EBITDA growth, potentially creating pressure to pursue more aggressive automation.

The central tension remained unresolved: if 100% of recommendations would soon originate from AI systems, as Gidwaney predicted, the human role would be reduced to orchestration and relationship management. Was that a sustainable competitive advantage or merely a transitional phase before full automation? The answer would depend on OneDigital's choices as well as how clients, competitors and the technology itself evolved.

Exhibit 1 OneDigital's AI Journey



Source: Company Documents

Exhibit 2 Five Dimensions of Workplace Transformation

HR Dimension	Traditional Practice	AI Disruption	Key Questions	Required Adjustments
Experience and Merit	Job descriptions specify years of experience and demonstrated skills. Merit-based advancement rewards employees for work they personally accomplished.	AI compresses learning into 'virtual apprenticeships.' Performance now depends on access to AI tools and skill in using them.	Does AI-assisted work represent the same professional development? Is it merit when advancement is based on work facilitated but not done by the person?	Revise job descriptions and hiring criteria. Redefine 'real-world' experience. Clarify merit-based promotion conditions.
Individual Contributors vs. Team Players	Employees categorized as individual contributors (work independently) or team players (collaborate with others).	Individual contributors working with AI are effectively in teams with non-human members. Performance depends on AI collaboration skills: prompting, evaluating outputs, integrating AI work.	How to assess social skills when many interactions occur with AI? Is someone who excels with AI but struggles with humans still a team player?	Redefine individual contributor when rarely alone. Develop frameworks for assessing AI collaboration alongside interpersonal skills.

HR Dimension	Traditional Practice	AI Disruption	Key Questions	Required Adjustments
Culture and Diversity	Culture transmitted through shared attitudes, traditions, and observation. Diversity focuses on promoting equal opportunities across backgrounds and perspectives.	AI systems must embody company values or risk cultural erosion. A terse AI undermines collaboration culture; efficiency-focused AI contradicts client-centric values.	How to embed culture into AI systems, not just document it? Does the AI skills gap matter more than traditional diversity metrics?	Design AI personality to reinforce culture. Consider AI skills gap as new inequality dimension. Expand definition of diverse thinking.
Performance and Development	Performance reviews evaluate what employee accomplished and how. Development through manager coaching, training programs, and career guidance.	AI contributes substantially to outcomes, making it difficult to separate individual from AI-assisted performance. AI provides real-time personalized coaching and hyper personalized learning plans.	Should reviews assess work products or skill in using AI to generate products? Should performance reviews evaluate employee and AI separately, then assess collaboration?	Create frameworks for evaluating AI-assisted performance. Balance algorithmic recommendations with holistic career guidance. Redefine manager role.
Leadership and Hierarchy	Leaders valued for judgment, ability to synthesize information, and skill in guiding teams. Authority and accountability flow through organizational hierarchy.	AI can analyze situations, generate recommendations, and provide individualized guidance. Can evaluate performance, inform resource allocation, support strategic decisions.	Is using AI for management decisions supportive leadership or delegation to algorithms? When AI can second-guess decisions, how do authority and accountability function?	Rethink organization structure and design. Clarify where accountability resides for AI-assisted decisions. Redefine leadership training priorities.

Source: Company Documents

Exhibit 3 AI Principles

Gift of Time Back

- We are a human capital organization – our “product” is the time that our people spend solving our client’s challenges
- Time management, mental bandwidth and burnout are the #1 inhibitors to our long-term success
- If our people have more time to spend with clients, our business will continue to grow and be successful
- **Our AI should be designed to give our people the gift of time back**

Better Insights

- The challenges we solve for our clients are increasingly complex in an ever-changing world
- Our clients expect us to know their business, know the industry and know the best solutions all the time
- Convergence is requiring our people to know more across different domains
- **Our AI should help our people have better insights into our business and our clients**


Source: Company Documents

Exhibit 4 Coworker: Map AI Features to Human Qualities

-  Our AI Coworkers behave and act like humans.
-  The more humanistic and HR-like we make these Coworkers, the more they will be adopted within our enterprise.
-  Our goal is to build a methodology that encompasses what it takes for us to “hire”, “onboard” and “manage” these AI Coworkers as they become effective members of our teams.
-  These Coworkers are at the beck-and-call of our humans, they are not autonomous, but they do have specific skill sets, memories and capabilities that superpower our humans.



Source: Company Documents

Exhibit 5 The Coworker Hiring Process

Source: Company Documents

Exhibit 6 AI Coworker Evaluation and Progression**Three Stage Career Progression**

Stage	Purpose	Exit Criteria	User Scope
<i>Intern</i>	Pilot Readiness	Must pass 4 of 5 criteria: - Core tasks validated in real use - Quality: satisfaction $\geq$ target - Reliability: no critical incidents - Evidence of value: measurable benefit - Adoption: meaningful sessions achieved	5-15 pilot users
<i>Apprentice</i>	Scale + Deeper Capability	Must pass 5 of 6 criteria: - Adoption: MAU $\geq$ target % of users - Engagement: interactions/user $\geq$ target - Quality: satisfaction $\geq$ target	20-100 users 2-3 weeks typical

		- Automation: % handled end-to-end - Measurable value: hours/cost saved - Risk/compliance: 0 incidents + audit complete	
Full-Time	Production + ROI Accountability	Must pass ALL 5 criteria: - Outcome KPIs trending toward targets - Operating KPIs stable (adoption not declining) - Trust KPI stable (quality not declining) - ROI story documented - Governance + escalation + access controls in place	Company-wide deployment

Performance Evaluation Scorecard


Category	Weight	Focus
Outcomes (Business KPIs)	40%	Measurable business value delivered <i>Examples: Hours saved, cost reduction, cycle time improvement, error reduction, revenue influenced</i>
Adoption/Engagement	20%	Usage patterns and user engagement <i>Examples: Monthly active users (MAU), interactions per user, task volume, automation rate</i>
Quality/Trust	25%	Output reliability and accuracy <i>Examples: User satisfaction %, QA accuracy rate, critical incidents (target: 0), compliance pass rate</i>
Ops Readiness	15%	Scalability and governance <i>Examples: Governance framework, QA plan, escalation path, compliance audit</i>

Decision Matrix

Score Range	Recommendation	Action
80-100	PROMOTE	Advance to next stage
60-79	KEEP	On track, continue monitoring
40-59	IMPROVE PLAN	Needs improvement before promotion
20-39	DEMOTE	Scale back deployment
0-19	RETIRE	Continue coworker

Source: Compiled from Company Documents

Exhibit 7Ben, the Benefits Expert



Source: Company Documents

Exhibit 8Examples of AI Coworkers

AI Coworkers


Growing roster across Practices



Coworker Name	Practice Area	Function
<i>Leigh</i>	Employee Benefits	Transforms complex renewal documents into focused internal summaries, professional client presentations, and concise email messages.
<i>Nina</i>	Employee Benefits	A Non-Qual insight analyst who dives deep into SEC Proxy Statements, extracts key details about Nonqualified Deferred Compensation (NQDC) plans, and produces strategic summaries and insights for outreach.
<i>Samara</i>	Employee Benefits	Pulls essential coverage details from Summary of Benefits Coverage (SBC) documents and organizes them precisely into a standard format, making benefits comparisons effortless.
<i>Ben</i>	Employee Benefits	Helps OneDigital's own employees understand the company's health and retirement plans.
<i>Max</i>	Employee Benefits	Provides platform support for Benemax
<i>ODysseus</i>	Finance/M&A	Strategic M&A research partner who helps identify, analyze, and evaluate potential acquisition targets while creating content for deal materials.
<i>Claus</i>	Finance/M&A	Reads an APA, checks it against OneDigital's standard template, and gives a clear list/table of what's different, missing, or added.
<i>Jackie</i>	Financial Services	A retirement 5500 Insight analyst who identifies insights, highlights risks and opportunities, and packages findings into consultant-ready reports.
<i>Avery</i>	Property & Casualty	A property researcher who quickly gathers comprehensive U.S. property data from multiple sources, saving hours of manual research.
<i>Quinn</i>	Property & Casualty	A Quote Comparison coworker who reads P&C PDF quotes and gives you clean comparison tables and quick summaries.
<i>Ace</i>	Cross-Practice	General virtual assistant for broad queries

Source: Compiled from Company Documents

Exhibit 9A Hybrid Workforce



Source: Company Documents